

Connecting to postgres

You should have

- a running postgresql 16 or 15 install on your machine
- pgAdmin installed

Server vs db

You connect to a postgres **server**

On that server you may have zero, one or multiple databases.

When you launch postgres on your local, you are actually launching the postgres server.

You connect and query the database on the command line using psql or in pgAdmin.

Users

The `postgres` user is the super user.

It can do anything on the server, create and manage databases, users, etc

Installing postgresQL also creates a user with your machine user name : *username*

The *username* user is not as powerful as the *postgres* user

Attention, on Mac, there is no *postgres* user. The superuser is your *username*

Connecting to the server

You can connect in the terminal or with pgAdmin

In the terminal you use `psql` to connect

psql takes the following parameters

- -U username
- -h host (IP address or local host)
- -p port
- -d database_name

PostgreSQL port is by default 5432 and does not have to be specified each time.

so to connect to a local instance of a postgres server with user bob and database treesdb

```
psql -U bob -d treesdb
```

to connect to a remote server IP: 123.456.789.246 same user and database

```
psql -U bob -h 123.456.789.246 -d treesdb
```

check your connection

to see your current connection run

```
psql -c "\conninfo"
```

if it complains of not knowing the database, add the *postgres* database as the database name

```
psql -c "\conninfo" -d postgres
```

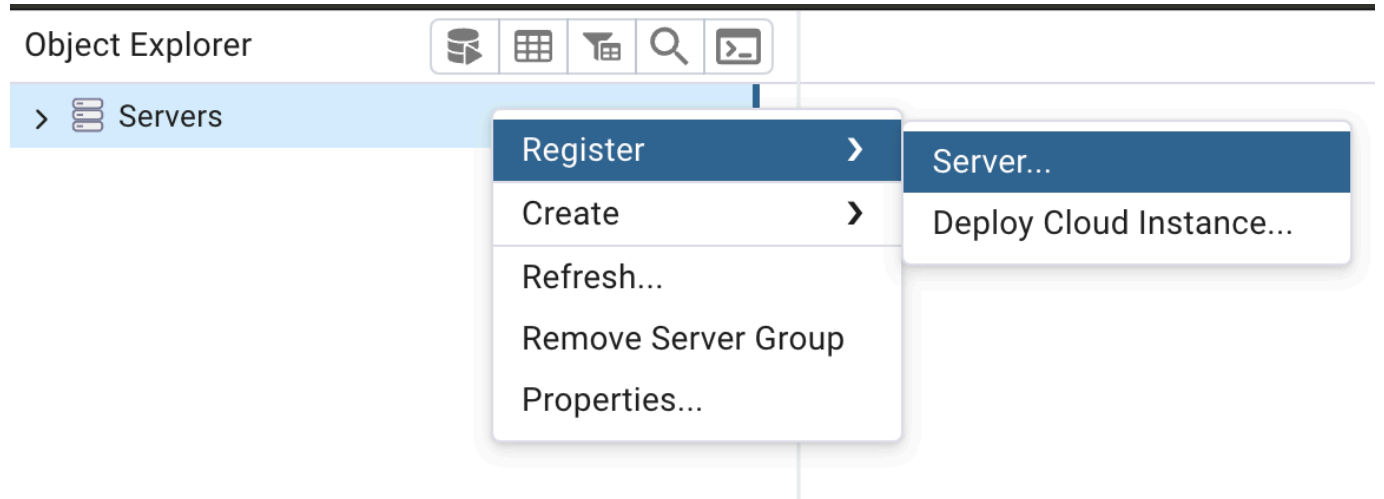
This should return:

```
You are connected to database "postgres" as user "alexis" via socket in "/tmp" at port "5432".
```

in pgAdmin,

Open pgAdmin

b. Right-click on "Servers" in the left panel and choose "Create" > "Server..." c. In the "Create - Server" dialog:



General tab:

Name: Give it a name like "Local PostgreSQL"

Connection tab:

- Host name/address: localhost
- Port: 5432 (default PostgreSQL port)
- Maintenance database: postgres
- Username: [your_username]
- Password: Leave blank if you haven't set one

local_server

General

Connection

Parameters

SSH Tunnel

Advanced

Host name/address

localhost

Port

5432

Maintenance database

postgres

Username

alexis

Kerberos authentication?

☐

Role

Service

?

?

Close

Reset

Save

to connect to a remote server, just change the connection tab parameters (host, user, password)

On windows

- Install postgres and follow the instructions
- click on the psql exe file
- it will open a terminal window and prompt you with hostname
 - if you're connecting on local just press enter
 - if you're connecting to remote add the IP address

Same with the other parameters

Loading a csv dataset

The steps are

1. create the database. let's call it treesdb. you can do that either
 1. on the command line with `createdb treesdb`
 2. in psql, with CREATE DATABASE ...
2. create the table let's call it `trees`
3. upload the data from the csv file into the trees table with

\COPY

Restore a database from a *sql dump*

Download the treesdb_v01.sql.gz file in the github repo

go to: <https://github.com/SkatAl/epitadb/tree/master/data> click right on the filename and `save link as`

You can also clone the repo if you have git installed

```
git clone git@github.com:SkatAI/epitadb.git
```

or just get the link to the file (click right on the filename and `copy link`) and use curl or wget to download the file

```
wget https://github.com/SkatAI/epitadb/blob/master/data/treesdb_v01.sql.gz
```

Once you have the file on your local you can restore it.

But first you must create the database either in pgAdmin (click right on server name and create > database; check the encoding is UTF8) or with the query

```
sql CREATE DATABASE treesdb_v01 WITH OWNER = alexis ENCODING = 'UTF8' LOCALE_PROVIDER = 'libc' CONNECTION LIMIT = -1 IS_TEMPLATE =
```

replace the owner name (alexis) with your username

The restore the database with

In the terminal

```
pgrestore --username "yourusername" --no-password --dbname "treesdb_v01" --section=pre-data --section=data --section=post-data --verbose "your path to/treesdbv01.sql.gz"
```

or in pgAdmin

In the restore dialog:

- Set "Format" to "Custom or tar"
- Browse and select your dump file (treesdb_v01.sql.gz).
- In the "Sections" tab, make sure "Pre-data", "Data", and "Post-data" are all checked. (not sure that's even required)

Check the data

You should have a single table called trees with the columns we saw last week.

```
treesdb_v01=# \d trees
```

Column	Type	Collation	Nullable	Default
idbase	integer			
location_type	character varying			
domain	character varying			
arrondissement	character varying			
suppl_address	character varying			
number	character varying			
address	character varying			
id_location	character varying			
name	character varying			
genre	character varying			
species	character varying			
variety	character varying			
circumference	integer			
height	integer			
stage	character varying			
remarkable	character varying			
geo_point_2d	character varying			

Accents and the Encoding

We need the database to be UTF8 encoded for the Paris trees data.

if you notice that the accent in some columns are not properly encoded for instance é is displayed as Ã©, à as Ã ...

```
UPDATE trees SET columnname = convertfrom(convertto(columnname, 'LATIN1'), 'UTF8')
```

do that for the columns : name, genre, species, variety, address ...

you can also check that server, client and table are UTF8 encoded with

```
SHOW serverencoding; SHOW clientencoding; SELECT pgencodingtochar(encoding) FROM pgdatabase WHERE datname = 'treesdb';
```

This is a version of the database with the proper encoding.

create the new database

- create a new database called treesdb_02
- set the user to your username (should appear in the dropdown)
- set the encoding to UTF8

The SQL tab should show

```
CREATE DATABASE treesdb_02
WITH
  OWNER = alexis
  ENCODING = 'UTF8'
  LOCALE_PROVIDER = 'libc'
  CONNECTION LIMIT = -1
  IS_TEMPLATE = False;
```

restore the database

Then click right on the treesdb_02 database and click on restore

- select the filename treesdb_v02.sql.gz,
- select format custom or tar
- click restore

The command line equivalent is

```
pg_restore --host "localhost" \
--port "5432" \
--username "alexis" \
--no-password \
--dbname "treesdb_02" \
"/Users/alexis/work/epitadb/data/treesdb_02.sql.gz"
```

while the db is restored, look at the process

Process Watcher - Restoring backup on the server

Restoring backup on the server 'local_server (localhost:5432)'

Running command:

```
/Applications/pgAdmin 4.app/Contents/SharedSupport/pg_restore --host "localhost" --port "5432" --
username "alexis" --no-password --dbname "treesdb_02" --verbose
"/Users/alexis/work/epitadb/data/treesdb_02.sql.gz"
```

🕒 Start time: Fri Sep 06 2024 09:20:31 GMT+0200 (Central European Summer Time)

⌛ Stop Process

```
pg_restore: connecting to database for restore
pg_restore: creating TABLE "public.trees"
pg_restore: processing data for table "public.trees"
```



Successfully completed.

Execution time: 0.87 seconds

Then check that the tree table has been created in the public schema of the treesdb_02 database

- select random rows
- \d trees